

Bird Classification System: Implementation Analysis

Jahnabi Roy (2022CS11094)
Abhinav Rajesh Shripad (2022CS11596)

1 Overview

This report details the implementation and training results of a PyTorch-based bird classification system. We tested our model with Train-Val-Test split of 80-10-10 and achieved an accuracy of 89.06% on the test dataset. The final model bird.pth was trained on Train-Val split of 80-20 and achieved a final Validation accuracy of 89.61%.

2 System Configuration

Learning Rate:	0.001	Number of Epochs:	80
Train-Val-Test Split:	80-10-10	Batch Size:	128
Early Stopping Patience:	20 epochs	Image Size:	64×64 pixels
Number of Classes:	10	Data Augmentation :	Ver/Hor/15° Rotate

3 Model Architecture

3.1 Neural Network Components

3.1.1 Background Removal Network

- **Encoder:**

- Input \rightarrow Conv2D(3 \rightarrow 64) \rightarrow BatchNorm2D \rightarrow ReLU \rightarrow MaxPool2D
- Conv2D(64 \rightarrow 128) \rightarrow BatchNorm2D \rightarrow ReLU \rightarrow MaxPool2D
- Conv2D(128 \rightarrow 256) \rightarrow BatchNorm2D \rightarrow ReLU \rightarrow MaxPool2D
- Conv2D(256 \rightarrow 512) \rightarrow BatchNorm2D \rightarrow ReLU \rightarrow MaxPool2D

- **Decoder:**

- ConvTranspose2D(512 \rightarrow 256) \rightarrow BatchNorm2D \rightarrow ReLU
- ConvTranspose2D(256 \rightarrow 128) \rightarrow BatchNorm2D \rightarrow ReLU
- ConvTranspose2D(128 \rightarrow 64) \rightarrow BatchNorm2D \rightarrow ReLU
- ConvTranspose2D(64 \rightarrow 1) \rightarrow Sigmoid

3.1.2 Classification Network


- Four convolutional blocks with structure:
 - Conv2D \rightarrow BatchNorm2D \rightarrow ReLU \rightarrow MaxPool2D
 - Channel progression: $3 \rightarrow 64 \rightarrow 128 \rightarrow 256 \rightarrow 512$
- Fully connected layers:
 - Flatten
 - Linear($8192 \rightarrow 1024$) \rightarrow BatchNorm1D \rightarrow ReLU
 - Linear($1024 \rightarrow 512$) \rightarrow BatchNorm1D \rightarrow ReLU
 - Linear($512 \rightarrow 10$)

3.2 Regularisation

- Early Stopping : 20 Epochs
- Batch Normalisation

4 Training Analysis

4.1 Training Progression



4.2 Training Observations

- Both Training and Validation Loss show a decreasing trend, though validation loss is not as smooth as the training loss. Similarly Training and Validation Accuracy increase as the training progresses. The progression is much smoother for training than for validation, however LR=0.001 output the best validation loss and accuracy than others.

5 Results

5.1 Regularisation

- With **Dropout** : **76%** accuracy on test set

- With **Batch Normalisation** : **80.5%** accuracy on test set
- With **Batch Normalisation, early stopping and adding data augmentation** : **89.06%** accuracy on test set

5.2 Performance Metrics

- Total Dataset Size: 9,960 images
- Train-Val-Test Split : 80-10-10
- Final Accuracy on Test Set: 89.06%

6 GradCAM Results

The GradCAM program was run over the validation dataset, which was also subject to transformations applied to train dataset (horizontal flip, vertical flip and at max 15 degrees rotation). Hence the images are pixellated. However the GradCAM results on the Classifier CNN do reveal what the model has learnt.




Figure-1 : Classes 0-3




Figure-2 : Classes 4-7




Figure-3 : Classes 8-9